

Car Price Prediction using RANDOM FOREST

In [1]:

```
import pandas as pd
```

Load Dataset from Local Directory

In [2]:

```
from google.colab import files
uploaded = files.upload()
```

Choose File

No file selected

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving Dataset.csv to Dataset.csv

Load Dataset

In [5]:

```
dataset = pd.read_csv('Dataset.csv')
dataset = dataset.drop(['car_ID'],axis = 1)
```

Summarize Dataset

In [6]:

```
print(dataset.shape)
print(dataset.head(5))
```

(205, 25)

	symboling	CarName	fueltype	...	citympg	highwaympg	price
0	3	alfa-romero giulia	gas	...	21	27	13495.0
1	3	alfa-romero stelvio	gas	...	21	27	16500.0
2	1	alfa-romero Quadrifoglio	gas	...	19	26	16500.0
3	2	audi 100 ls	gas	...	24	30	13950.0
4	2	audi 100ls	gas	...	18	22	17450.0

[5 rows x 25 columns]

Splitting Dataset into X & Y

This X contains Numerical & Text Dataset

In [7]:

```
Xdata = dataset.drop('price', axis = 'columns')
numericalCols = Xdata.select_dtypes(exclude = ['object']).columns
X = Xdata[numericalCols]
X
```

Out[7]:

	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower
0	3	88.6	168.8	64.1	48.8	2548	130	3.47	2.68		9.0
1	3	88.6	168.8	64.1	48.8	2548	130	3.47	2.68		9.0
2	1	94.5	171.2	65.5	52.4	2823	152	2.68	3.47		9.0
3	2	98.0	176.6	66.2	54.2	3227	160	2.40	2.40		10.0

symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower
4	2	99.4	176.6	66.4	54.3	2824	136	3.19	3.40	8.0
...	...	...	...	...	...	...	...	...	...	...
200	-1	109.1	188.8	68.9	55.5	2952	141	3.78	3.15	9.5
201	-1	109.1	188.8	68.8	55.5	3049	141	3.78	3.15	8.7
202	-1	109.1	188.8	68.9	55.5	3012	173	3.58	2.87	8.8
203	-1	109.1	188.8	68.9	55.5	3217	145	3.01	3.40	23.0
204	-1	109.1	188.8	68.9	55.5	3062	141	3.78	3.15	9.5

205 rows x 14 columns



In [8]:

```
Y = dataset['price']
Y
```

Out[8]:

```
0      13495.0
1      16500.0
2      16500.0
3      13950.0
4      17450.0
...
200     16845.0
201     19045.0
202     21485.0
203     22470.0
204     22625.0
Name: price, Length: 205, dtype: float64
```

Scaling the Independent Variables (Features)

In [11]:

```
from sklearn.preprocessing import scale
cols = X.columns
X = pd.DataFrame(scale(X))
X.columns = cols
X
```

Out[11]:

	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower
0	1.743470	-1.690772	-0.426521	0.844782	-2.020417	-0.014566	0.074449	0.519071	1.839377	-0.288349	
1	1.743470	-1.690772	-0.426521	0.844782	-2.020417	-0.014566	0.074449	0.519071	1.839377	-0.288349	
2	0.133509	-0.708596	-0.231513	0.190566	-0.543527	0.514882	0.604046	2.404880	0.685946	-0.288349	
3	0.938490	0.173698	0.207256	0.136542	0.235942	-0.420797	-0.431076	0.517266	0.462183	-0.035973	
4	0.938490	0.107110	0.207256	0.230001	0.235942	0.516807	0.218885	0.517266	0.462183	-0.540725	
...	...	...	...	...	...	...	...	...	...	...	...
200	-1.476452	1.721873	1.198549	1.398245	0.728239	0.763241	0.339248	1.666445	0.336970	-0.162161	
201	-1.476452	1.721873	1.198549	1.351515	0.728239	0.949992	0.339248	1.666445	0.336970	-0.364062	
202	-1.476452	1.721873	1.198549	1.398245	0.728239	0.878757	1.109571	0.926204	1.232021	-0.338824	

	symboling	wheelbase	carlength	carwidth	carheight	curbweight	engineize	boreratio	stroke	compressionratio	horsepower
203	-1.476452	1.721873	1.198549	1.398245	0.728239	1.273437	0.435538	1.183483	0.462183	3.244916	
204	-1.476452	1.721873	1.198549	1.398245	0.728239	0.975021	0.339248	1.666445	-	-0.162161	

205 rows x 14 columns

## Splitting Dataset into Train & Test

In [12]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size = 0.20, random_state = 0)
```

## Traing using RANDOM FOREST

In [13]:

```
from sklearn.ensemble import RandomForestRegressor
model = RandomForestRegressor()
model.fit(x_train, y_train)
```

Out[13]:

```
RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                      max_depth=None, max_features='auto', max_leaf_nodes=None,
                      max_samples=None, min_impurity_decrease=0.0,
                      min_impurity_split=None, min_samples_leaf=1,
                      min_samples_split=2, min_weight_fraction_leaf=0.0,
                      n_estimators=100, n_jobs=None, oob_score=False,
                      random_state=None, verbose=0, warm_start=False)
```

## Evaluating Model

In [15]:

```
ypred = model.predict(x_test)

from sklearn.metrics import r2_score
r2score = r2_score(y_test, ypred)
print("R2Score", r2score*100)
```

R2Score 90.67976965632458